Exploratory Data Analysis on Housing Market in Italy

Supervised Learning & Visualization

Daniel Anadria Kyuri Park Ernst-Paul Swens Emilia Löscher

October 3, 2022

Contents

1 Introduction

This is an exploratory data analysis of the Italian housing market in 2022. For context, Italy contains a total of 20 regions (*regioni*), 107 provinces (*province*) and 7,904 municipalities (*comuni*). In the present work, we pose several interesting research questions which can be answered by means of data visualization and summary statistics.

1.1 The Dataset

Our dataset originates from Kaggle. It contains information about the housing market in Italy in 2022. The data were scraped from one of the most prominent housing sales websites in Italy during the month of *August 2022*. The data consist of more than 223,000 sales posts spread over 7,023 (89% coverage) Italian municipalities. We do not have any information on the representativeness of our dataset. Hence, we advise caution when drawing inferences from our findings.

In order to plot the statistics of interest to maps of Italy, we use the regional and provincial shape files, which are obtained from the *Italian National Institute of Statistics* (ISTAT). These files contain the regional and provincial coding and geographical shape information, which can be used to cluster the municipalities in our `location` variable into their respective provinces and regions.

For each housing sale post, the dataset contains the following variables:

Table 1: Description of Variables in the Italy Housing Dataset

Variable	**Description**
`id`	ID of the sale
`timestamp`	Timestamp consisting of 10 digits
`location`	Location on municipality level
`title`	Short description of property
`price`	Price in Euros
`n_rooms`	Number of rooms
`floor`	Floor
`mq`	Size in square meters
`n_bathrooms`	Number of bathrooms
`year_of_construction`	Year of construction
`availability`	Availability of property
`energy_class`	Energy class ranging from a+ to g
`status`	Status of the property
`heating`	Type of heating
`has_garage`	Garage present: yes (1), no (0)
`has_terrace`	Terrace present: yes (1), no (0)
`has_garden`	Garden present: yes (1), no (0)
`has_balcony`	Balcony present: yes (1), no (0)
`has_fireplace`	Fireplace present: yes (1), no (0)
`has_alarm`	Alarm present: yes (1), no (0)
`has_air_conditioning`	Air Conditioning present: yes (1), no (0)
`has_pool`	Pool present: yes (1), no (0)
`has_parking`	Parking present: yes (1), no (0)
`has_elevator`	Elevator present: yes (1), no (0)
`is_furnished`	Furniture present: yes (1), no (0)

2 Preparation

In order to start our exploratory analysis, we first load relevant packages and import the dataset as well as the ISTAT shape files.

2.1 Load Packages & Import Data

```
## load packages
library(tidyverse)   # for wrangling data
library(skimr)       # for skimming data
library(sf)          # for spatial analysis
library(sp)          # for spatial analysis
library(ggplot2)     # for plotting
library(fuzzyjoin)   # for joining on not-exact matches
library(ggpubr)      # for arranging ggplots
library(mice)        # for imputation procedure

## import italy housing data
houses <- read.csv("data/housing_data_italy_august2022.csv",
                   na.strings=c("","NA"), header = TRUE)

## import istat shape files
# municipality
municipalities <- st_read("data/italy_shape_2022_files/Com01012022_g")
municipalities <- municipalities[c("COD_REG", "COD_PROV", "COMUNE")]
# province
provinces <- st_read("data/italy_shape_2022_files/ProvCM01012022_g")
# region
regions <- st_read("data/italy_shape_2022_files/Reg01012022_g")
```

3 Exploratory Research Questions

The goal of our analysis is:

1. Exploring the missingness in the dataset.
2. Investigating whether there are any geographical trends in the median housing prices and their absolute deviations on regional and/or provincial level?

4 Data Cleaning

Note: We base the following data cleaning on the summary of the raw dataset, which can be found in the *Appendix*.

The original data consist of 223,409 rows (*sales*) and 25 columns (*variables*).
Given our research questions, we exclude `id` (*ID of the sale*), `timestamp` (*timestamp of the sale*), and

`title` (*description of the property*) as they are deemed irrelevant. In addition, we exclude two columns that have only one unique value (`status` and `availibility`), as these are not variables but constants.

We observe that the types of some variables are wrongly specified. We convert them to a correct type (e.g., `heating`: character→factor, `has_xxx`: numeric→factor, `is_furnished`: numeric→factor).

Next, we create a new variable `property_age` by subtracting the `year_of_construction` from 2022. In the original dataset, there are some unreasonable years of construction (e.g., 2209). While some properties may be sold before their construction is completed, we deem it unlikely for properties whose `year_of_construction` is more than 4 years later as of now. Thus, we filter out those with `year_of_construction` $>$ 2026.

For the second exploratory research question, our variable of interest is `price`. We notice that it is highly skewed to the right given that the mean (239,939) is far off to the right of the median (135,000). We examine the distribution of `price` with a boxplot to examine the outliers (see *Figure 1*).

```
## create our own theme that can be used throughout
custom.theme = theme(
  axis.title.x = element_text(size = 14),
  axis.text.x = element_text(size = 13),
  axis.title.y = element_text(size = 14),
  axis.text.y = element_text(size = 13))

## boxplot of price
houses %>%
   ggplot(aes(x=price)) +
   geom_boxplot() +
   # add comma on the x-axis labels
   scale_x_continuous(labels=scales::label_comma(),
   # rotate the x-axis labels
   guide = guide_axis(angle = 25)) +
   # apply minimal theme plus our own theme
   theme_minimal() + custom.theme
```

From *Figure 1*, we observe that there are extreme outliers in `price`. Some housing prices in the dataset are exorbitant (e.g., over €2B). We decide to focus the scope of our analysis on the houses whose price is less than or equal to €1M. The distribution of housing prices after filtering can be seen in *Figure 2*.

```
## density plot of price (cleaned dataset)
houses %>%
  # filter the price over a million
  filter(price <= 1e6 | is.na(price)) %>%
  # create a ggplot
  ggplot(aes(price)) +
  # add histogram
  geom_histogram(aes(y=..density..), bins = 30, color = 1, fill="white") +
  # add density line
  geom_density(lwd=0.5, color = "#165e70", fill = "#165e70", alpha = 0.2) +
  # apply our theme
  theme_minimal() + custom.theme
```

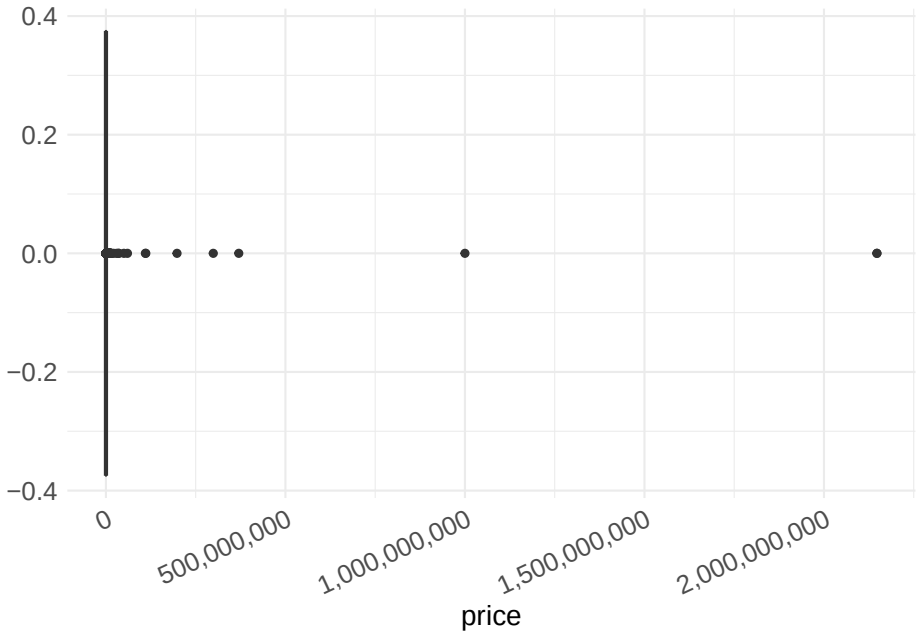

Figure 1: Boxplot of Housing Prices in Italy

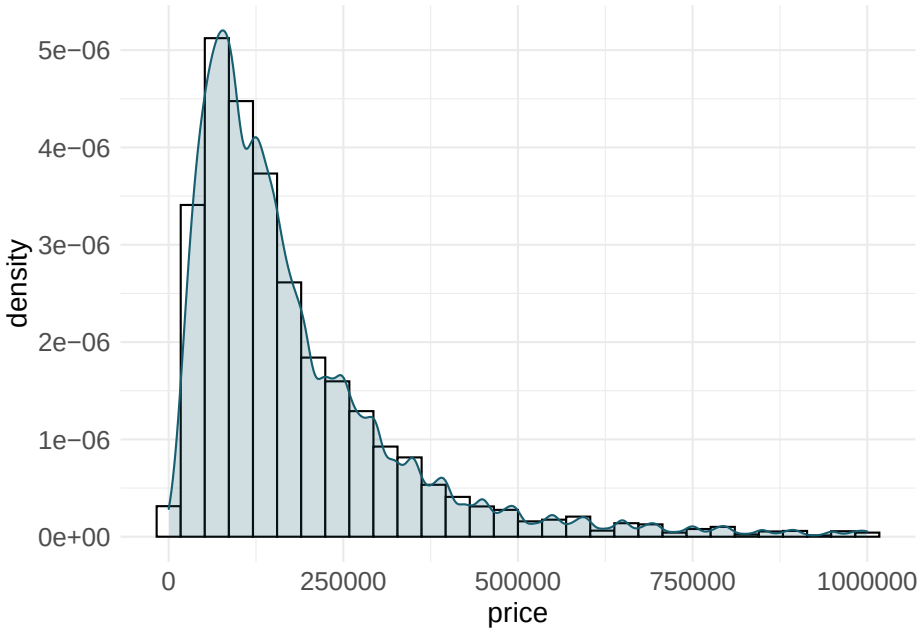

Figure 2: Histogram and Density Plot of Housing Price After Filtering

From *Figure 2*, we take that the distribution of housing prices after filtering shows that there is still a long right tail, but that is to be expected with housing prices in any country. The extreme outliers have been eliminated. From this plot, we also conclude that when working with housing price data, it is advisable to use centrality and spread measures that are robust to skewed data. For this reason, we will use the *median* and *median absolute deviation (MAD)* instead of the mean and variance in our exploration of the present dataset.

```
## cleaning up the housing data
houses.cleaned <- houses %>%
  # only select variables that have more than one unique value
  select(where(~n_distinct(.) > 1)) %>%
  # fix the data type (convert them to factor)
  mutate(across(c(starts_with("has"), is_furnished, heating, energy_class, n_rooms,
                  n_bathrooms, location), factor)) %>%
  # filter out houses whose price is over a million (while keeping NAs)
  filter(price <= 1e6 | is.na(price),
  # filter out houses with a built year over 2026 (while keeping NAs)
         year_of_construction < 2026 | is.na(year_of_construction)) %>%
  # create property age variable
  mutate(property_age = 2022 - as.numeric(year_of_construction)) %>%
  # remove id, timestamp, title and year_of_construction
  select(-c(id, timestamp, title, year_of_construction))
```

As some of our explorations are on the regional and provincial, we use the *ISTAT* shape files to append our original dataset with provincial and regional information. To this end, we use fuzzy string matching for inexact matches as we found that there were some minor inconsistencies in how the municipalities were named in our dataset as opposed to their names in the ISTAT shape files.

```
data.geo <- stringdist_left_join(houses.cleaned, municipalities,
         by = c("location" = "COMUNE"), distance_col = "distance",
         ignore_case = TRUE) %>%
      group_by(location) %>%
      slice_min(distance) %>%
      select(-geometry,-distance) %>%
      left_join(., as.data.frame(regions[,c("COD_REG","DEN_REG")])) %>%
      select(-geometry, -COMUNE) %>%
      left_join(., as.data.frame(provinces[,c("DEN_UTS", "COD_PROV")],
         by = "COD_PROV")) %>%
      rename(region = DEN_REG, province = DEN_UTS) %>%
      select(-geometry, -COD_REG, -COD_PROV) %>%
      relocate(c(region, province), .after = location)
```

4.1 Data Summary

After data cleaning, we take a look at the summary statistics to get a better overview of our data. We skim through our cleaned dataset using the `skimr` package.

```r
## settings: round up by 2 decimal places & disable scientific notation
options(digits = 3, scipen = 999)
## specify skimming function
custom.skim <- skim_with(character = sfl(whitespace=NULL, min = NULL, max = NULL),
   factor = sfl(ordered = NULL),
   numeric = sfl(p0 = NULL, p25=NULL, p50=NULL, p75=NULL, p100=NULL, hist=NULL,
                 median = ~median(., na.rm=T),
                 min = ~min(., na.rm=T), max = ~max(., na.rm=T), n_unique=n_unique))
## summary table
custom.skim(houses.cleaned)
```

Table 2: Data summary

Name	houses.cleaned
Number of rows	220748
Number of columns	20
_____	
Column type frequency:	
factor	16
numeric	4
_____	
Group variables	None

Variable type: factor

skim_variable	n_missing	complete_rate	n_unique	top_counts
location	0	1.00	7023	pis: 192, leg: 190, la : 188, bar: 187
n_rooms	58297	0.74	4	3: 56766, 4: 47144, 5: 30944, 2: 27597
n_bathrooms	12677	0.94	3	1: 107372, 2: 79843, 3: 20856
energy_class	638	1.00	12	g: 115238, f: 25396, e: 17124, a: 15931
heating	0	1.00	2	aut: 197849, oth: 22899
has_garage	0	1.00	2	0: 180669, 1: 40079
has_terrace	0	1.00	2	0: 196111, 1: 24637
has_garden	0	1.00	2	0: 184426, 1: 36322
has_balcony	0	1.00	2	0: 198058, 1: 22690
has_fireplace	0	1.00	2	0: 208817, 1: 11931
has_alarm	0	1.00	2	0: 218752, 1: 1996
has_air_conditioning	0	1.00	2	0: 155058, 1: 65690
has_pool	0	1.00	2	0: 216473, 1: 4275
has_parking	0	1.00	2	0: 217364, 1: 3384
has_elevator	0	1.00	2	0: 208067, 1: 12681
is_furnished	0	1.00	2	0: 203644, 1: 17104

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	median	min	max	n_unique
price	39113	0.82	177784.71	153812.51	130000	1	1000000	2555
floor	71398	0.68	1.82	1.13	2	1	52	20
mq	3343	0.98	156.08	124.68	116	1	999	971
property_age	10	1.00	56.07	74.36	42	-3	1022	375

From the output, we see that our cleaned dataset has 220,607 rows and 20 columns, 16 of which are factors, and 4 of which are numeric types. The output is presented in a table for factor and numeric variables, respectively.

From the table of factor variables, we again see that `location` has 7023 unique values (i.e., municipalities). Also, there are some missing values for `energy_class` and a lot of missing values for `n_rooms` and `n_bathrooms`. The ratio of properties having an alarm is highest (99.10%) and lowest for having air conditioning (70.24%). Three rooms, one bathroom, autonomous heating and energy class g were the most frequent levels of the respective factors.

In the numeric table, we see that all the variables have some missing values. About 18% of `price` is missing. We look into this more in detail when we address the first exploratory question which concerns the missingness in `price`. Furthermore, we see that the median `floor` is 2 and the maximum is 52. The average size in square meters (`mq`) for this dataset is 156.1 m^2 and while the oldest building according to the dataset is 1022 years old, the average age of the buildings is 56.1 years.

5 Answering the Exploratory Questions

5.1 Question 1: Exploring missingness in the dataset.

5.1.1 Visual Exploration

First, we calculate and plot the proportion of missing values for all the variables in *Figure 3 A*.

```
## barplot to show the missingness per variable
houses.cleaned %>%
   # get the average missing proportion
   is.na() %>% colMeans() %>% stack() %>%
   # create ggplot for the missing proportion
   ggplot(aes(x = reorder(ind, values), y=values)) +
   # specify the bar plot
   geom_bar(stat="identity", width = .2, fill = "black",
            position = position_dodge2(padding = 0.8)) +
   # flip the bar plot and add the points
   geom_point() + coord_flip() +
   # apply our custom theme
   theme_minimal() + custom.theme +
   # specify axis labels
   labs(x = "", y = "Missingness (%)")
```

The plot shows that the percentage of missing values lies above 30% for the variable *floor* and above 25% for the *number of rooms*. Furthermore, about 18% of the *housing prices* are missing. Except for the

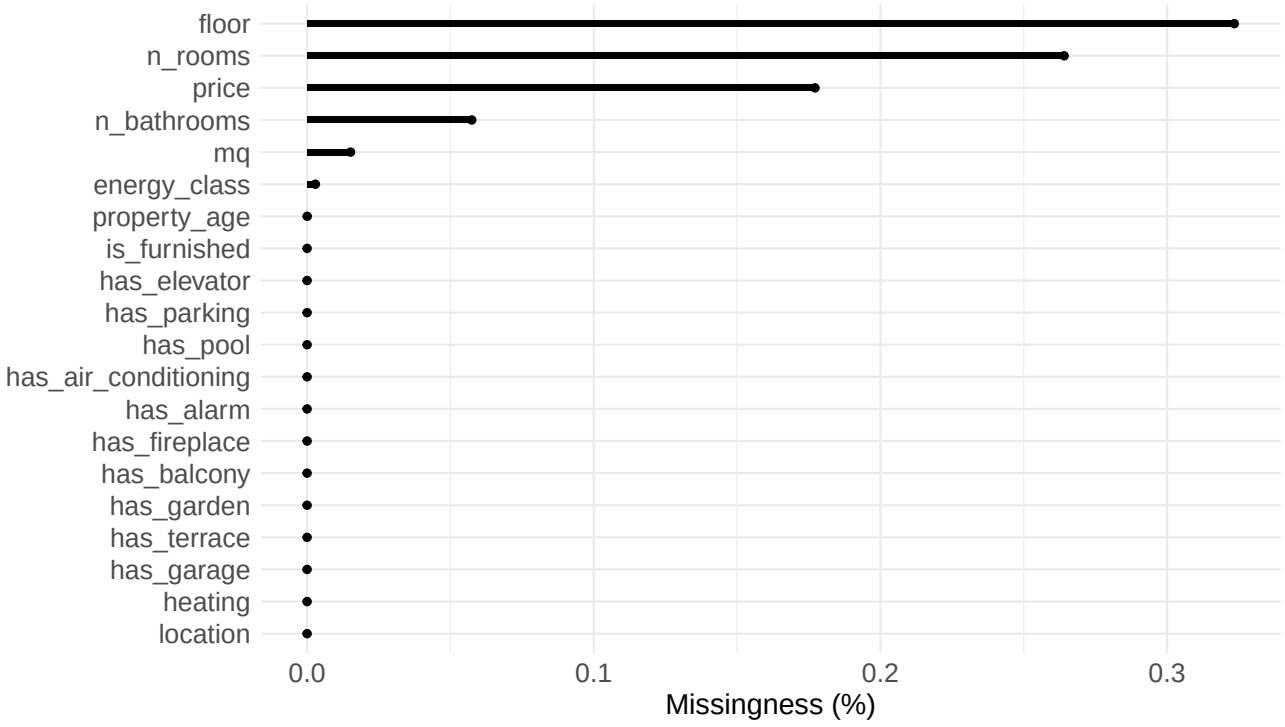

Figure 3: Proportion of Missing Values

number of bathrooms which has about 6% values missing, for the remaining variables less than 3% are missing, respectively.

Next we explore the correlation of *price* and other variables and summarize the findings in *Table 5*.

```
houses.cleaned %>%
  # create missingness indicator of price
  mutate(na_ind = ifelse(is.na(price), 1, 0)) %>%
  ungroup() %>%
  # compute correlation
  summarize(across(n_rooms:property_age, ~cor(as.numeric(.x), na_ind,
                                              use = "pairwise.complete.obs"))) %>%
  # sort them in descending order based on their absolute values
  t(.) %>% .[order(abs(.), decreasing = TRUE),] %>%
  # create a neat table
  knitr::kable(col.names = "correlation", caption = "Correlation between missingness
               of price and other variables", align="c")
```

Table 5: Correlation between missingness of price and other variables

	correlation
energy_class	-0.272
has_garden	-0.114
heating	0.109

	correlation
has_garage	-0.097
has_fireplace	-0.087
has_air_conditioning	-0.086
has_terrace	-0.078
has_elevator	-0.077
n_rooms	0.062
has_balcony	-0.056
has_parking	-0.038
mq	0.033
has_alarm	-0.032
property_age	-0.020
n_bathrooms	0.017
floor	0.010
has_pool	-0.010
is_furnished	0.005

The missingness of `price` appears to be moderately correlated with the `energy_class` (*cor = -0.271*). Hence, we further check what the pattern of missingness in price across different energy classes looks like.

```
## create plots for missingness in price vs energy class
houses.cleaned %>%
  # add price missingness indicator
  mutate(na_ind = ifelse(is.na(price), 1, 0)) %>%
  # group by energy class
  group_by(energy_class) %>%
  # sum up all the missingness in price per energy class
  summarize(n = sum(na_ind)) %>%
  # create a ggplot for the sum of missingness
  ggplot(aes(x = energy_class, y = n, fill = energy_class)) +
  # turn off legend
  geom_col(show.legend=FALSE) +
  # customize the color
  scale_fill_manual(values = c("#F8766D", "#999999","#F8766D",
                               rep("#999999", 8),"#F8766D")) +
  # apply the themes to the bar plot
  theme_minimal() + custom.theme +
  # change the labels
  labs(x = "Energy Class", y = "Missingness in price")
```

Figure 4 shows that there is higher missingness in price for houses that either have a good energy class of *a* or *a1*, or fall into the very inefficient energy class *g* (marked in red). There is low missingness in *price* for the other energy classes and for houses for which *energy_class* is missing. ***@Kyuri "most colour blind people are unable to fully 'see' red, green or blue light." Fig. 4 might use another color***

We also check the missingness in price across different regions to see if there are any geographical patterns.

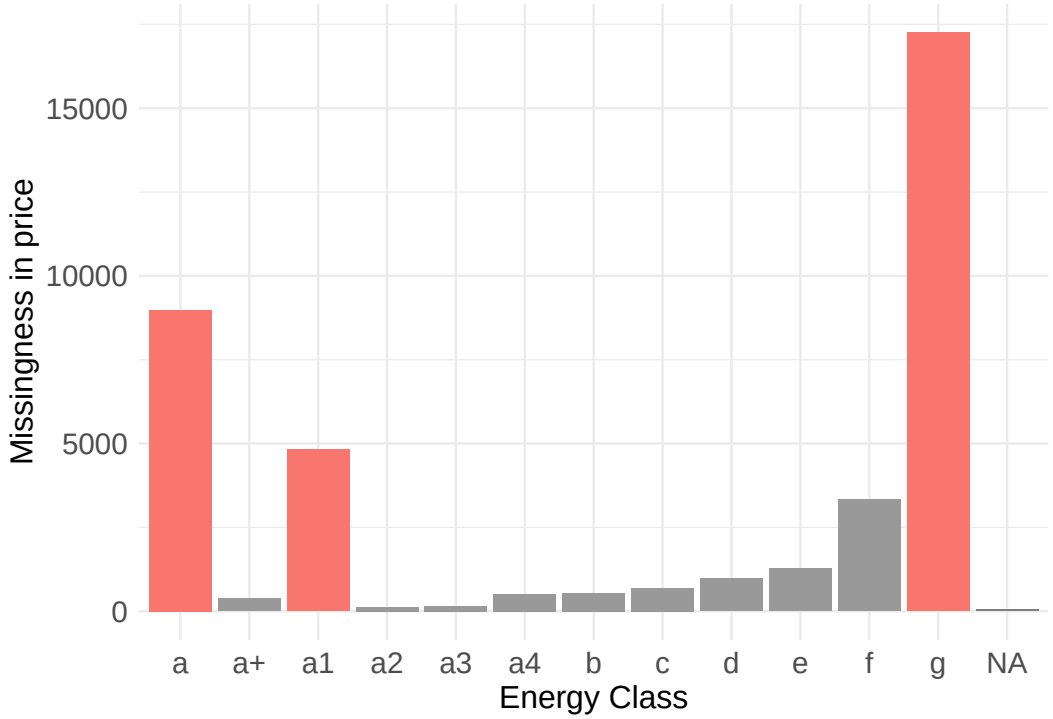

Figure 4: Missing Values in Price across Different Energy Classes

```
## check missingness in price w.r.t regions
data.geo %>%
   # add price missingness indicator
   mutate(na_ind = ifelse(is.na(price), 1, 0)) %>%
   group_by(region) %>%
   # get the total missingness proportion per region
   summarize(`Average missing proportion (%)` = sum(na_ind) / n()) %>%
   # add the spatial data
   left_join(., regions, by = c("region" = "DEN_REG")) %>%
   st_as_sf() %>%
   ggplot() +
   # plot the italy map
   geom_sf(fill=NA) +
  # add scatter plots of missingness proportion per region
   geom_point(color = alpha("red", 0.4),
   aes(size = `Average missing proportion (%)`, geometry = geometry),
   stat = "sf_coordinates") +
  scale_size(range = c(1, 5)) +
  # remove unnecessary coordinates
   theme_void() +

   theme(legend.position = "bottom")
```

Figure 5 shows that the proportion of missing data in `price` is not equal across different regions. However, there seem to not be any regions which are drastically more likely to have missing price data/

Figure 5: Missingness of Price per Region

Next, we examine the relationship between the predictors' missingness indicators and the *price* densities by plotting the density of price (on a log10-scale) split by whether or not the respective value for a given variable is missing or not. These results provide us with some indication of the missing data mechanism. Specifically, if missing completely at random (MCAR) is plausible to assume or not.

```
## function to create a density plot of price with missing indicator
missing.plots <- function(x) {
   plot <- houses.cleaned %>% ungroup() %>%
      mutate(missing = is.na(.[,x])) %>%
      ggplot(aes(x = log10(price), fill = missing)) +
      geom_density(alpha = 0.5, color = NA) +
      theme_classic() +
      labs(x = expression(paste(Log["10"],"(Price)"), y = "Density")) +
      scale_x_continuous(limits = c(2, 7)) + ylim(0, 1.5)
   return(plot)
}

## multiple plots for floor, no. of rooms, no. of bathrooms, and meters squared
ggarrange(missing.plots("floor"),
          missing.plots("n_rooms"),
          missing.plots("n_bathrooms"),
          missing.plots("mq"),
          labels = c("         A. Floor", "B. Number of rooms", "C. Number of bathrooms", "D. Squ
                     font.label=list(color="black",size = 10),
          ncol = 2, nrow = 2, common.legend = TRUE, legend = "bottom")
```

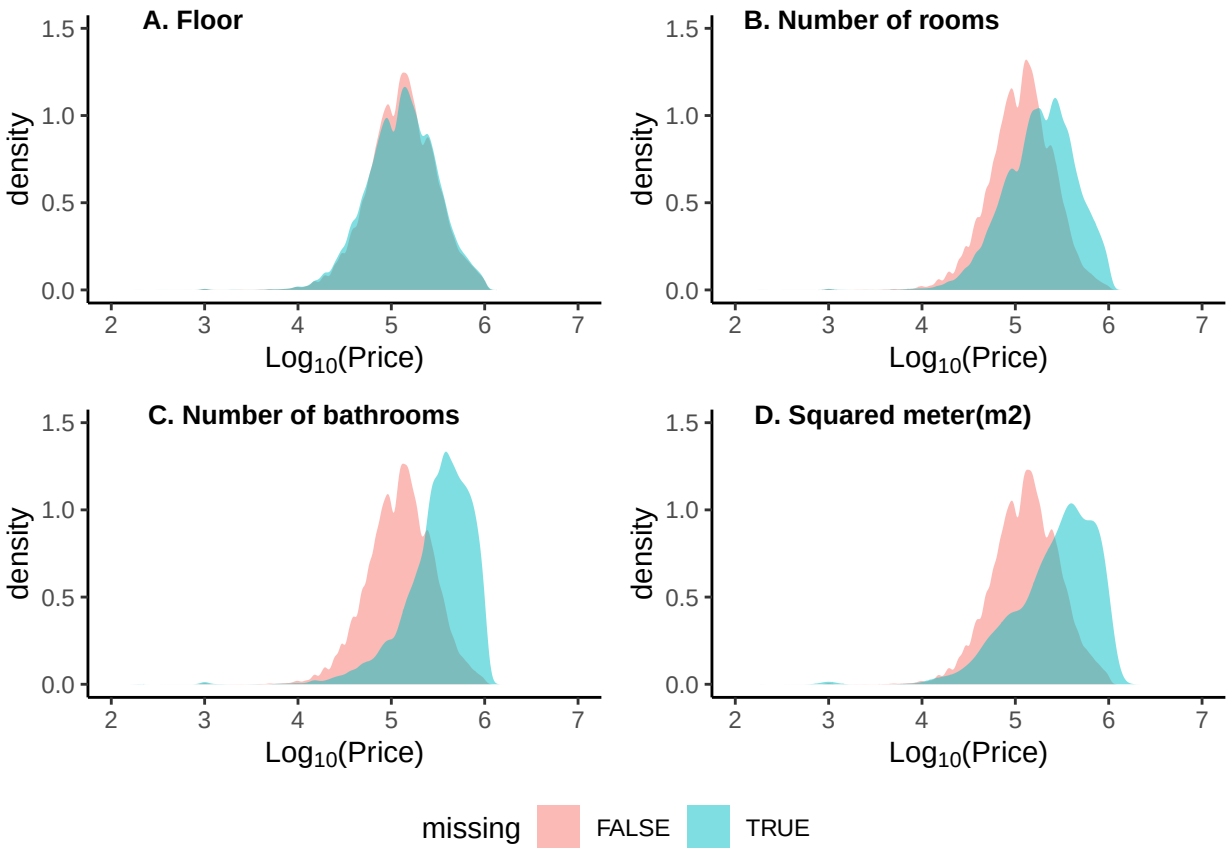

Figure 6: Density Plots for Price based on the Missing Indicators

The missingness of number of floors (*Figure 6 A*) and number of rooms (*Figure 6 B*) does not seem to be dependent on the observed price information. Whereas for the number of bathrooms (*Figure 6 C*) and the meters squared (*Figure 6 D*) missingness shows a different result, namely, missingness tends to occur at higher house prices. We conclude that the missingness mechanism is not MCAR. Therefore, we shouldn't use case-wise deletion. Instead, imputation is an appropriate technique to deal with our missing values.

5.1.2 Imputation

We run five multiple imputations on our data using a selection of predictors derived from `quickpred`. Note that some of the following code chunks are set to `eval=FALSE`. They take a long time to run and need not be rerun as we saved the resulting objects and load them in subsequent chunks.

@EP could you 1) add save.RDS lines at the bottom of the eval = F chunks and 2) write a little bit more on what is being done here?

```
## create predictor matrix
predictors <- houses.cleaned %>% ungroup() %>% quickpred()

## multiple imputation procedure (m = 5)
imputations <- mice(houses.cleaned, m = 5, seed = 1, predictorMatrix = predictors)
```

We limited the diagnostics for the imputation procedure to the convergence of the algorithm and plausibility of the imputed data. There appeared to be no convergence issues and the imputed data appeared to be plausible with respect to the observed data (see Appendix).

```
## load imputation results
imputations <- readRDS("data/imputation.rds")
```

```
## function to join geospatial data
join.geo <- function(data) {
   data.geo <- stringdist_left_join(data, municipalities,
         by = c("location" = "COMUNE"), distance_col = "distance",
         ignore_case = TRUE) %>%
      group_by(location) %>%
      slice_min(distance) %>%
      select(-geometry,-distance) %>%
      left_join(., as.data.frame(regions[,c("COD_REG","DEN_REG")])) %>%
      select(-geometry, -COMUNE) %>%
      left_join(., as.data.frame(provinces[,c("DEN_UTS", "COD_PROV")],
         by = "COD_PROV")) %>%
      rename(region = DEN_REG, province = DEN_UTS) %>%
      select(-geometry, -COD_REG, -COD_PROV) %>%
      relocate(c(region, province), .after = location)
}

## join geospatial data to the imputed data
imputed <- lapply(complete(imputations, "all"), join.geo)

## join geospatial data to the cleaned house data
houses.cleaned <- join.geo(houses.cleaned)
```

```r
## load the imputed data including the geospatial data
imputed <- readRDS("data/imputed.rds")

## load the cleaned house data including the geospatial data
houses.cleaned <- readRDS("data/housescleaned.rds")
```

5.1.3 Conclusion

For the first question, we explored missingness in the dataset. We saw that there are quite a lot of missing values for the *floor*, the *number of rooms*, the *number of bathrooms*, and *house price*. As *house price* is of interest for the second question, we took a closer look and realized that missingness on *price* is correlated, e.g., with *energy class* and that the amount of missingness varies across regions.

We also examined if the missingness of the variable *house price* and the other variables are related. From our results it appears that they are. Hence, it is unlikely that the missingness mechanism is MCAR. Therefore, we include a multiple imputation procedure.

5.2 Question 2: Regional and Provincial Trends in the Median Housing Price and the Median Absolute Deviations in Italy

We use the imputed dataset to explore whether there are geographical trends in the median and the median absolute deviation (MAD) *housing price*.

First we plot the median and the MAD on the regional and the provincial level:

```r
## function to get aggregated region price information
group.region <- function(data, estimate) {
   data %>%
      group_by(region) %>%
      summarize(estimate = ifelse(estimate, median(price), mad(price))) %>%
      select(-region)
}

region.median <- lapply(imputed, group.region, estimate = T) %>%
   do.call(cbind, .) %>% rowMeans()
region.mad <- lapply(imputed, group.region, estimate = F) %>%
   do.call(cbind, .) %>% rowMeans()

labels <- c("Trentino-Alto Adige", "Molise", "Piemonte", "Bolzano", "Calabria")

price.by.region <- data.frame(region = sort(regions$DEN_REG),
   median = region.median, mad = region.mad) %>%
   left_join(.,regions, by = c("region" = "DEN_REG")) %>% st_as_sf() %>%
   mutate(label = ifelse(region %in% labels, region, NA))
```

```r
## function to get aggregated province price information
group.province <- function(data, estimate) {
   data %>%
      group_by(province) %>%
```

```
      summarize(estimate = ifelse(estimate, median(price), mad(price))) %>%
      select(-province)
}

province.median <- lapply(imputed, group.province, estimate = T) %>%
   do.call(cbind, .) %>% rowMeans()
province.mad <- lapply(imputed, group.province, estimate = F) %>%
   do.call(cbind, .) %>% rowMeans()

price.by.province <- data.frame(province = sort(provinces$DEN_UTS),
   median = province.median, mad = province.mad) %>%
   left_join(., provinces, by = c("province" = "DEN_PROV")) %>% st_as_sf()
```

```
plot.list.1 <- list()

## median & mad of price per region
plot.list.1 <- map(
  c("median", "mad"),
  function(var) {
    ggplot(price.by.region) +
      # map each statistic
      geom_sf(aes(fill = .data[[var]])) +
      # void theme: remove all unncessary coordinates
      theme_void() +
      # add labels to specific regions
      geom_sf_label(aes(label = label)) +
      # color-scheme (color-blind friendly???)
      scale_fill_viridis_c(option = "E", direction = -1) +
      # lengthen the legend
      theme(legend.key.width= unit(2, 'cm'))
     }
  )
plot.list.2 <- list()

## median & mad of price per province
plot.list.2 <- map(
  c("median", "mad"),
  function(var) {
    ggplot(price.by.province) +
      # map each statistic
      geom_sf(aes(fill = .data[[var]])) +
      # void theme: remove all unncessary coordinates
      theme_void() +
      # color-scheme (color-blind friendly???)
      scale_fill_viridis_c(option = "E", direction = -1) +
      # lengthen the legend
      theme(legend.key.width= unit(2, 'cm'))
     }
  )
```

```
## combine the plot lists
plot.list <- c(plot.list.1, plot.list.2)
```

```
## plot the median of price
ggarrange(plotlist = plot.list[c(1,3)], nrow = 1, ncol = 2, common.legend = TRUE,
          legend = "bottom")
```

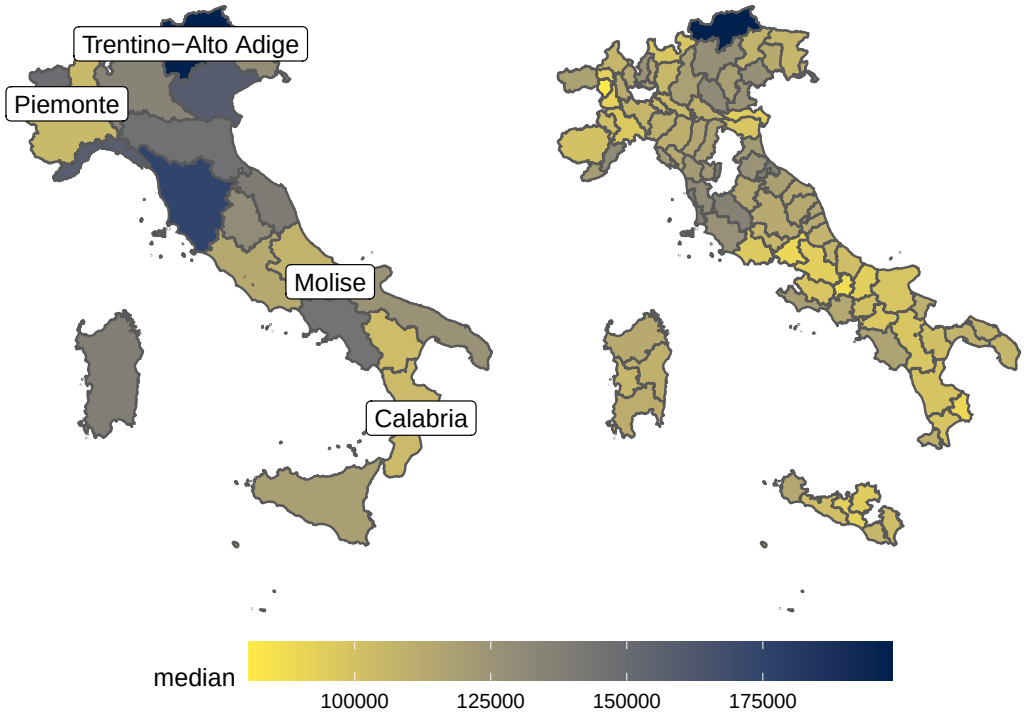

Figure 7: Median of Price per Region (left) and Province (right)

```
## plot the mad of price
ggarrange(plotlist = plot.list[c(2,4)], nrow = 1, ncol = 2, common.legend = TRUE,
          legend = "bottom")
```

On the regional level, we see that the median for the region *Trentino-Alto Adige* (€200,000) are the highest, and the lowest for *Molise* (€79,000). We recognize a trend that the median price is lower for more Southern regions in Italy. The only exception from this is the region of *Piemonte* (€99,000) which has a lower median price than the surrounding regions in the North. Regarding the MAD, a measure of variability within a region, we see that it is highest for the regions with higher median prices. This is recognizable as the color patterns in the median plot and the MAD plot are very similar, and it means that more expensive regions also tend to have more variability than less expensive regions.

On the provincial level, it can be seen that the high median of the *Trentino-Alto Adige* region is mainly due to the high median of the province of *Bolzano* (€400,000). As other provinces in that region have lower median prices, the MAD in that region is comparably high. The opposite is true within the southern regions where the provinces that make up the region of *Calabria* for example, all have a low median price. Hence, the MAD for that region is low.

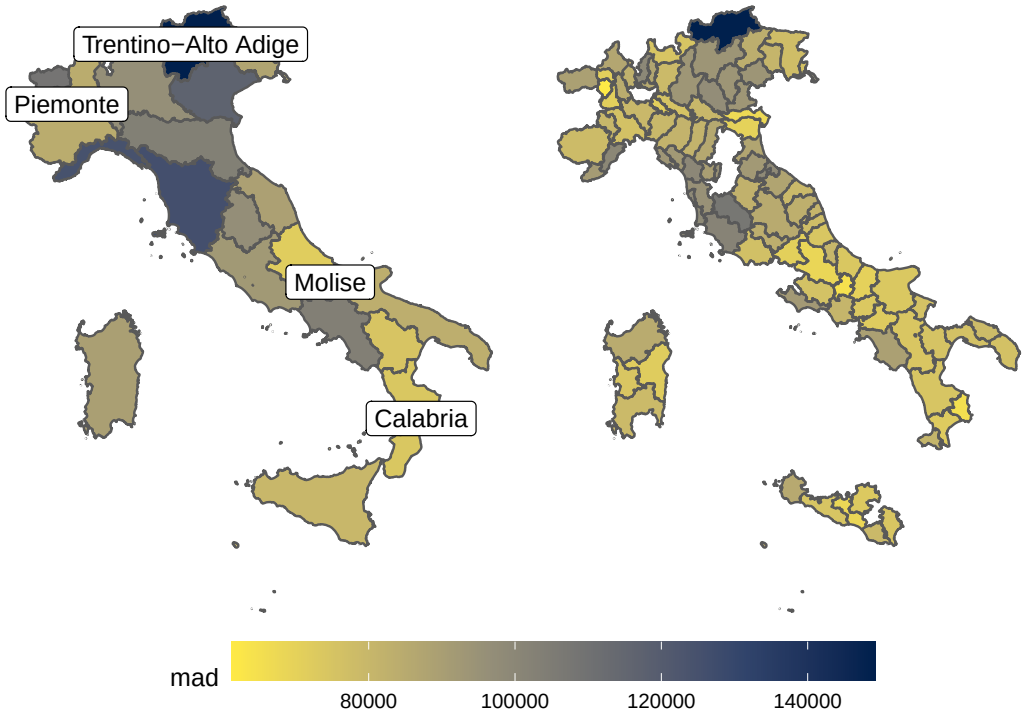

Figure 8: MAD of Price per Region (left) and Province (right)

Given that the overall geographical pattern for median and MAD of `price` correspond to each other, it is interesting to investigate further the possible differences in the distribution of `price` between high- and low-median regions.

```
## top 2 high median regions
top.2.median <- price.by.region %>% slice_max(median, n = 2) %>% pull(region)

## lowest 2 low median regions
lowest.2.median <- price.by.region %>% slice_min(median, n = 2) %>% pull(region)

## plot the histograms for each region in the high and low median groups
imputed[[1]] %>%
  # subset top two and bottom two countries
  filter(region %in% c(top.2.median, lowest.2.median)) %>%
  # group by the regions
  group_by(region) %>%
  # create the grouping variable for coloring
  mutate(grouping = ifelse(region %in% top.2.median, "high median", "low median"),
         # get the median price for each region
         med_price = median(price, na.rm=T)) %>%
  # create ggplot for price (coloring by groups)
  ggplot(aes(x = price, fill = grouping)) +
  geom_histogram(bins = 30) +
  # create a panel of plots per region
  theme_bw() + facet_wrap(~region) +
  # indicate the median price by a vertical line
```

```
geom_vline(aes(xintercept = med_price, group=region), linetype="dashed") +
# change legend title
labs(fill = "high/low regions")
```

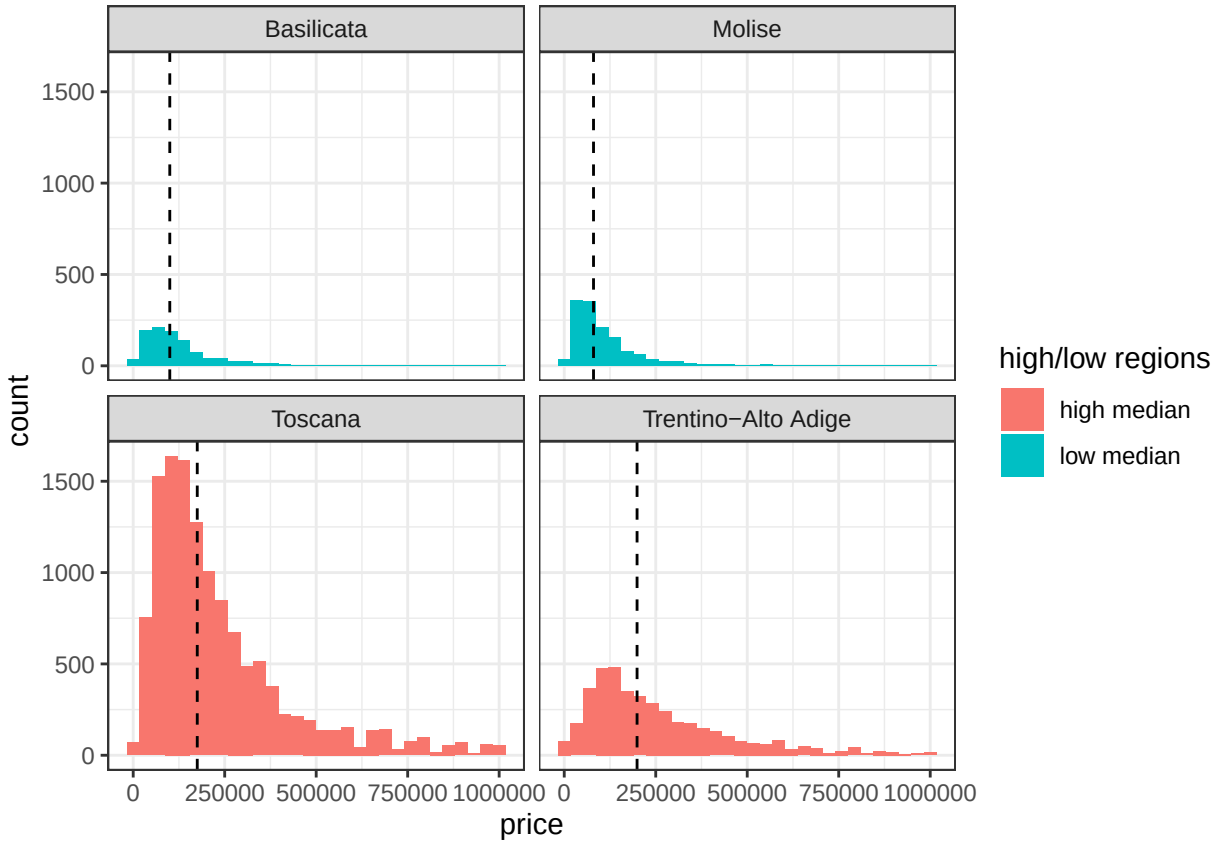

Figure 9: group differences in price distribution

An interesting take-away from *Figure 9* is that all the densities for the *house price* are right skewed. This is reasonable for house prices as one would expect that there are more cheap and moderately priced houses and only few very expensive houses. Furthermore, it is apparent that there are most sells in the dataset from the *Toscana* region and there are only very few very expensive houses in the regions of *Calabria* and *Molise* as one would expect with those regions having a lower median. ***maybe make this more clear***

5.2.1 Conclusion

All things considered, we found a general decrease in the median price when moving from northern to southern Italy. We also showed that higher medians on a regional level are due to a higher median on the province level with one region having a particular high median. Consequently, such regions where there is one province with a higher median of *house price*, the MAD is higher too if the other provinces in that region do not have a high median. Furthermore, by including the plots on the province level, it was recognizable that there were no sales in some provinces. ***How many provinces did not have data?***

6 Overall Conclusion

All things considered, we identified that there is quite some missingness present - especially for *house price* for which about 18% of values were missing. The absolute correlation of missingness in *house price* was highest with *energy class* (-0.27). As we were not able to conclude from our analysis that the missingness mechanism is MCAR, we decided to impute the missing values.

Using the imputed dataset, we explored the median and the MAD on the regional and provincial level in Italy. We found a trend in the median price from more expensive to cheaper when going from North to South in Italy. We also have shown that these higher medians on a regional level are due to a higher median on the provincial level. In most cases, the MAD for such regions was high, too, as there frequently is one province with a higher median of *house price* and the MAD is consequently higher if the other provinces in that region do not have a high median, too.

7 Appendix

7.1 Summary of the raw data using the my_skim function.

```
custom.skim(houses)
```

Table 6: Data summary

Name	houses
Number of rows	223409
Number of columns	25
_____	
Column type frequency:	
character	6
numeric	19
_____	
Group variables	None

Variable type: character

skim_variable	n_missing	complete_rate	empty	n_unique
location	0	1	0	7023
title	0	1	0	199305
availability	0	1	0	1
energy_class	679	1	0	12
status	0	1	0	1
heating	0	1	0	2

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	median	min	max	n_unique
id	0	1.00	111705.00	64492.77	111705	1	223409	223409
timestamp	0	1.00	1661135705.37	72645.42	1661135577	1661114079	1661158618	42238
price	39116	0.82	239938.98	7562062.01	135000	1	2147483647	2852
n_rooms	60323	0.73	3.50	0.99	3	2	5	4
floor	72365	0.68	1.82	1.13	2	1	52	22
mq	4034	0.98	158.63	128.68	117	1	999	976
n_bathrooms	14397	0.94	1.59	0.67	1	1	3	3
year_of_construction	10	1.00	1965.13	76.75	1980	1000	2209	389
has_garage	0	1.00	0.18	0.38	0	0	1	2
has_terrace	0	1.00	0.11	0.32	0	0	1	2
has_garden	0	1.00	0.17	0.37	0	0	1	2
has_balcony	0	1.00	0.10	0.30	0	0	1	2
has_fireplace	0	1.00	0.05	0.23	0	0	1	2
has_alarm	0	1.00	0.01	0.10	0	0	1	2
has_air_conditioning	0	1.00	0.30	0.46	0	0	1	2
has_pool	0	1.00	0.02	0.15	0	0	1	2
has_parking	0	1.00	0.02	0.12	0	0	1	2
has_elevator	0	1.00	0.06	0.23	0	0	1	2
is_furnished	0	1.00	0.08	0.27	0	0	1	2

7.2 Convergence of the algorithm and plausability of the imputed dataset

```
## convergence of the algorithm
plot(imputations)
```

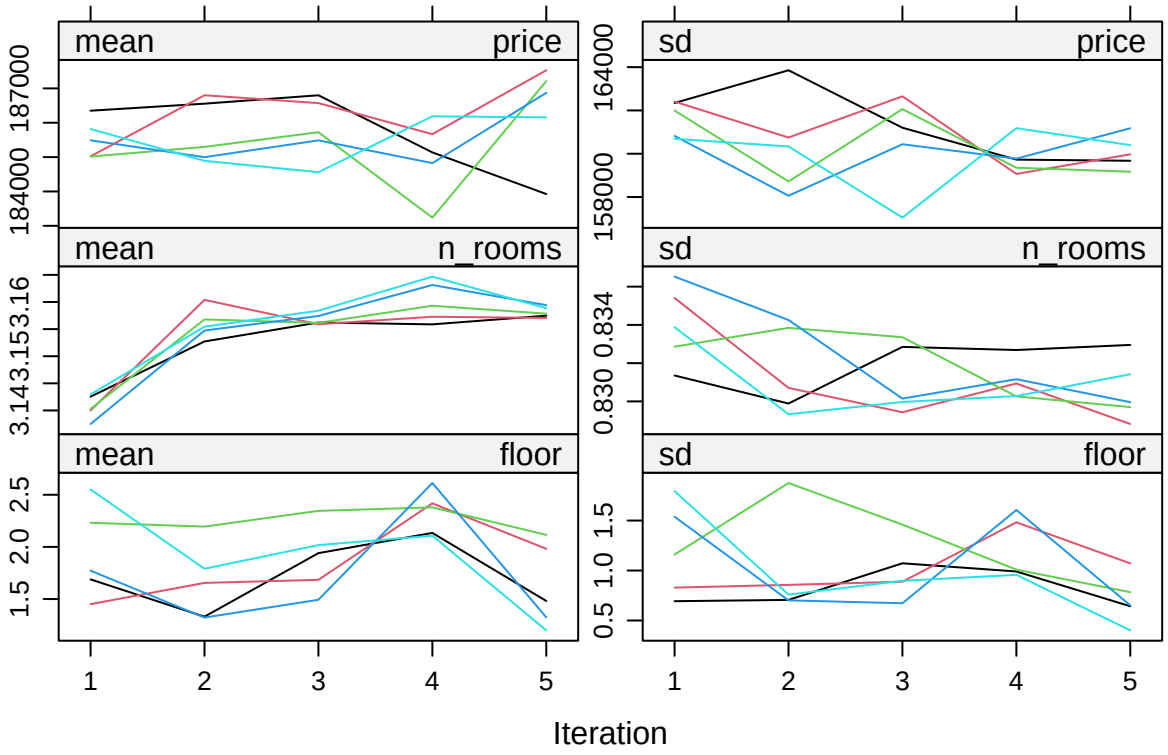
mean
price
sd
price
mean
n_rooms
sd
n_rooms
mean
floor
sd
floor
Iteration

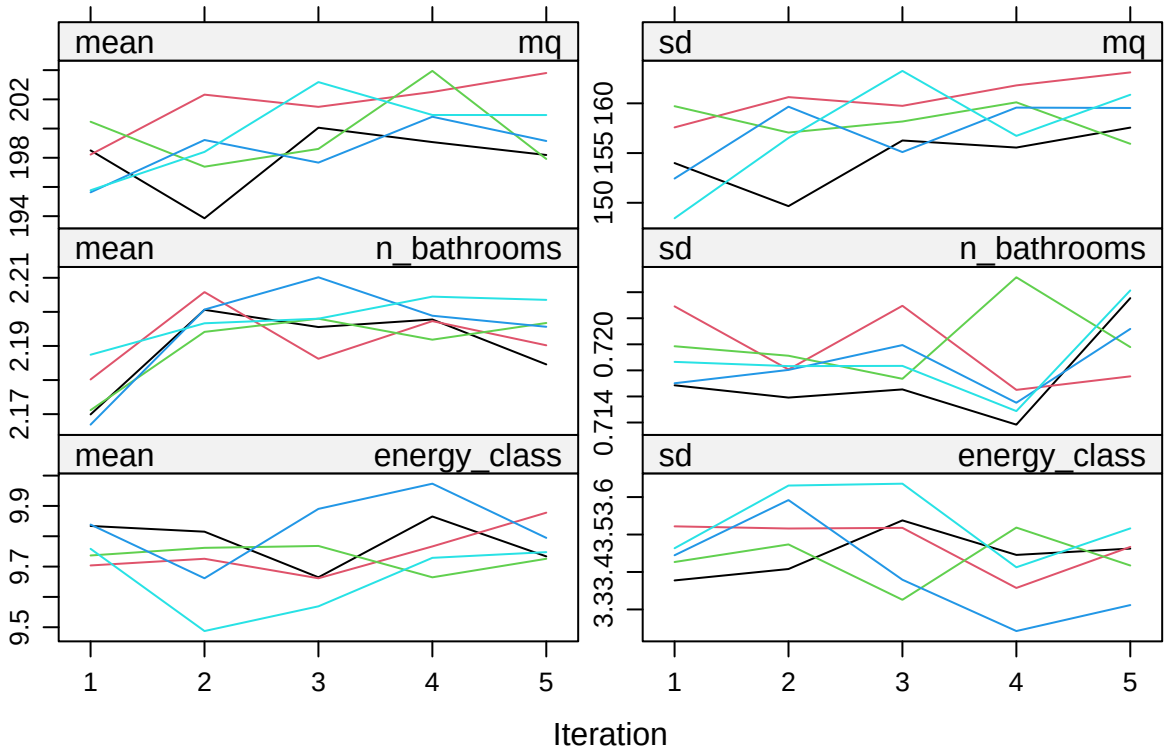
mean
mq
sd
mq
mean
n_bathrooms
sd
n_bathrooms
mean
energy_class
sd
energy_class
Iteration

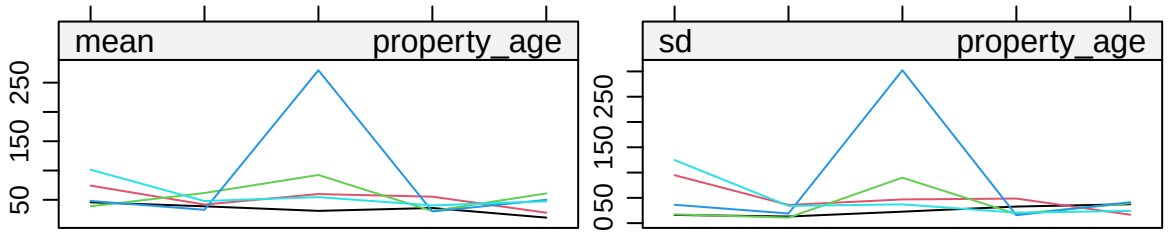

Iteration

```
## plausibility of the imputed data
densityplot(imputations, ~n_rooms + mq + floor + n_bathrooms + price, lwd = 2)
```

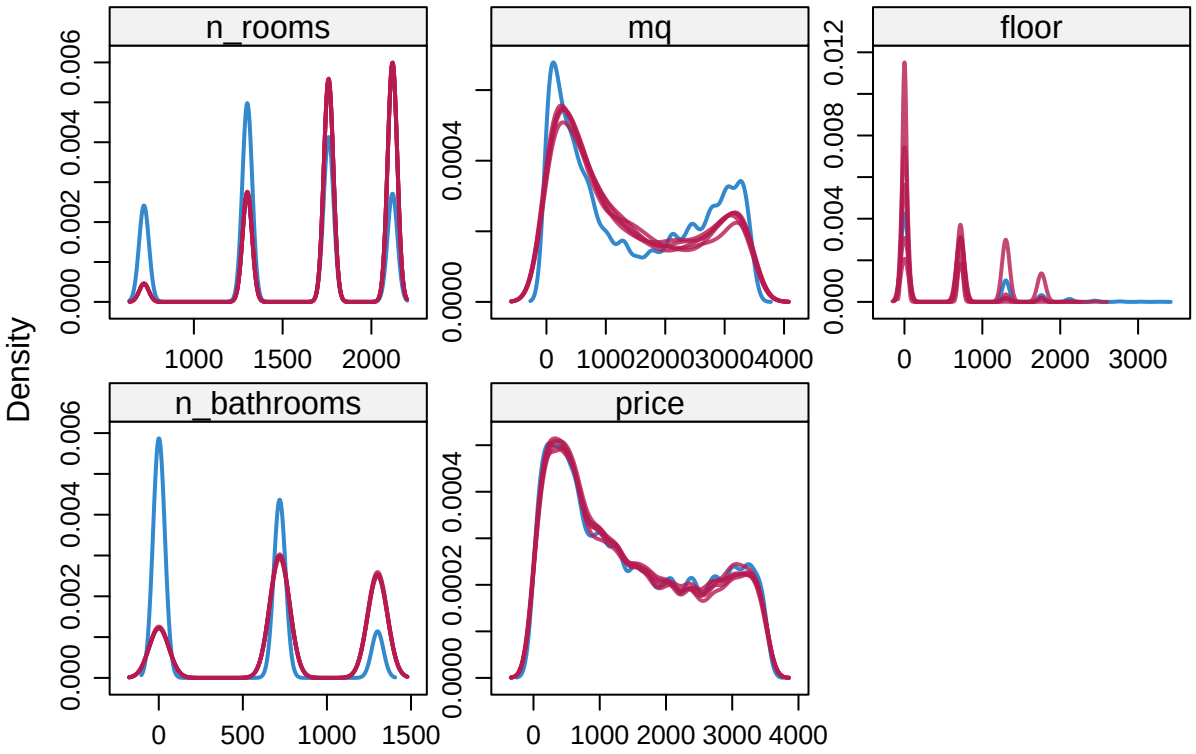
Density
n_rooms
mq
floor
n_bathrooms
price
0.000
0.002
0.004
0.006
0.0000
0.0004
0.000
0.004
0.008
0.012
1000
1500
2000
0
1000
2000
3000
4000
0
1000
2000
3000
0.0000
0.0002
0.0004
0
500
1000
1500